

## Phase 3 - Modeling

**Note 1:** the following starting code only generates a single random train/test split when `default_seed` is used. You need to modify the code to generate 100 independent train/test splits with different seeds and report the average results on those independent splits along with standard deviation.

**Note 2:** You are completely free to use your own implementation.

```
#mount google drive
from google.colab import drive
drive.mount('/gdrive')
```

```
#check files
!ls -la '/gdrive/My Drive/Case Study MLPS/'
```

```
📁 Drive already mounted at /gdrive; to attempt to forcibly remount, call drive.mount("/g
total 823153
-rw----- 1 root root 267687592 Apr 26 17:18 clean_data.pickle
-rw----- 1 root root 1208381 Apr 26 17:18 'CS-Phase 2.ipynb'
-rw----- 1 root root 610800 May 3 12:32 'CS-Phase 3.ipynb'
-rw----- 1 root root 355 Apr 26 17:19 'MLPS Phase 3.ipynb'
-rw----- 1 root root 1 May 1 20:20 'Phase 3 Write-up.gdoc'
-rw----- 1 root root 1 Apr 12 21:05 'Phase II Write-Up.gdoc'
-rw----- 1 root root 573399301 Apr 29 23:16 pre_clean_data.pickle
```

```
# Load general utilities
# -----
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.axes as ax
import datetime
import numpy as np
import pickle
import time
import seaborn as sns

# Load sklearn utilities
# -----
from sklearn.model_selection import train_test_split
from sklearn import preprocessing
from sklearn.model_selection import GridSearchCV

from sklearn.metrics import accuracy_score, classification_report, roc_auc_score, roc_curve, l

from sklearn.calibration import calibration_curve

# Load classifiers
# -----
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import RidgeClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import BaggingClassifier

# Other Packages
# -----
from scipy.stats import kendalltau
```

```

from sklearn.neural_network import MLPRegressor
from sklearn import linear_model
from sklearn.ensemble import RandomForestRegressor
from sklearn.cluster import KMeans
from sklearn.externals.six import StringIO
from IPython.display import Image
from sklearn.tree import export_graphviz
from scipy.interpolate import spline

# Load debugger, if required
#import pixiedust
pd.options.mode.chained_assignment = None # 'warn'

# suppress all warnings
import warnings
warnings.filterwarnings("ignore")

# Define a function that, given a CVGridSearch object, finds the
# percentage difference between the best and worst scores
def find_score_variation(cv_model):
    all_scores = cv_model.cv_results_['mean_test_score']
    return( np.abs((max(all_scores) - min(all_scores))) * 100 / max(all_scores) )

    ...

    which_min_score = np.argmin(all_scores)

    all_perc_diff = []

    try:
        all_perc_diff.append( np.abs(all_scores[which_min_score - 1] - all_scores[which_min_score]) )
    except:
        pass

    try:
        all_perc_diff.append( np.abs(all_scores[which_min_score + 1] - all_scores[which_min_score]) )
    except:
        pass

    return ( np.mean(all_perc_diff) )
    ...

# Define a function that checks, given a CVGridSearch object,
# whether the optimal parameters lie on the edge of the search
# grid
def find_opt_params_on_edge(cv_model):
    out = False

    for i in cv_model.param_grid:
        if cv_model.best_params_[i] in [ cv_model.param_grid[i][0], cv_model.param_grid[i][-1] ]:
            out = True
            break

    return out

```

## ▼ Define a default random seed and an output file

```

default_seed = 1
output_file = "output_sample"

# Create a function to print a line to our output file

def dump_to_output(key, value):

```

```
with open(output_file, "a") as f:
    f.write(",".join([str(default_seed), key, str(value)]) + "\n")
```

## ▼ Load the data and engineer the features

```
# Read the data and features from the pickle file saved in CS-Phase 2
data, discrete_features, continuous_features, ret_cols = pickle.load( open( "/gdrive/My Drive,
```

```
data.head(5)
```

	id	loan_amnt	funded_amnt	term	int_rate	installment	grade	emp_length	l
0	1077501	5000.0	5000.0	36 months	10.65	162.87	B	10+ years	
1	1077430	2500.0	2500.0	60 months	15.27	59.83	C	< 1 year	
2	1077175	2400.0	2400.0	36 months	15.96	84.33	C	10+ years	
3	1076863	10000.0	10000.0	36 months	13.49	339.31	C	10+ years	
4	1075358	3000.0	3000.0	60 months	12.69	67.79	B	1 year	

5 rows × 32 columns

```
## Create the outcome columns: True if loan_status is either Charged Off or Default, False otherwise
data["outcome"] = np.where(data['loan_status'].isin(['Charged Off', 'Default']), True, False)
```

```
# Create a feature for the length of a person's credit history at the time the loan is issued
continuous_features = list(continuous_features)
data['cr_hist'] = (data.issue_d - data.earliest_cr_line) / np.timedelta64(1, 'M')
continuous_features.append('cr_hist')
```

```
# Randomly assign each row to a training and test set. We do this now because we will be fitting
np.random.seed(default_seed)
## create the train columns where the value is True if it is a train instance and False otherwise
data['train'] = np.random.choice([True, False], len(data), p=[0.7, 0.3])
```

```
# Create a matrix of features and outcomes, with dummies. Record the names of the dummies for
X_continuous = data[continuous_features].values
```

```
X_discrete = pd.get_dummies(data[discrete_features], dummy_na = True, prefix_sep = ":", drop_prefix = True)
discrete_features_dummies = X_discrete.columns.tolist()
X_discrete = X_discrete.values
```

```
X = np.concatenate( (X_continuous, X_discrete), axis = 1 )
```

```
y = data.outcome.values
```

```
train = data.train.values
```

## ▼ Prepare functions to fit and evaluate models

```
# see 3.1.2. in the PDF for an explanation of the split
def prepare_data(data_subset = np.array([True]*len(data)),
                 n_samples_train = 30000,
                 n_samples_test = 20000,
                 feature_subset = None,
                 date_range_train = (data.issue_d.min(), data.issue_d.max()),
                 date_range_test = (data.issue_d.min(), data.issue_d.max()),
                 random_state = default_seed):
    ...

This function will prepare the data for classification or regression.
It expects the following parameters:
- data_subset: a numpy array with as many entries as rows in the
  dataset. Each entry should be True if that row
  should be used, or False if it should be ignored
- n_samples_train: the total number of samples to be used for training.
  Will trigger an error if this number is larger than
  the number of rows available after all filters have
  been applied
- n_samples_test: as above for testing
- feature_subset: A list containing the names of the features to be
  used in the model. In None, all features in X are
  used
- date_range_train: a tuple containing two dates. All rows with loans
  issued outside of these two dates will be ignored in
  training
- date_range_test: as above for testing
- random_state: the random seed to use when selecting a subset of rows

Note that this function assumes the data has a "Train" column, and will
select all training rows from the rows with "True" in that column, and all
the testing rows from those with a "False" in that column.

This function returns a dictionary with the following entries
- X_train: the matrix of training data
- y_train: the array of training labels
- train_set: a Boolean vector with as many entries as rows in the data
  that denotes the rows that were used in the train set
- X_test: the matrix of testing data
- y_test: the array of testing labels
- test_set: a Boolean vector with as many entries as rows in the data
  that denotes the rows that were used in the test set
...

np.random.seed(random_state)

# Filter down the data to the required date range, and downsample
# as required
filter_train = ( train & (data.issue_d >= date_range_train[0]) &
                 (data.issue_d <= date_range_train[1]) & data_subset ).values
filter_test = ( (train == False) & (data.issue_d >= date_range_test[0])
               & (data.issue_d <= date_range_test[1]) & data_subset ).values

filter_train[ np.random.choice( np.where(filter_train)[0], size = filter_train.sum()
                                - n_samples_train, replace = False ) ] = False
filter_test[ np.random.choice( np.where(filter_test)[0], size = filter_test.sum()
                               - n_samples_test, replace = False ) ] = False

# Prepare the training and test set
X_train = X[ filter_train , :]
X_test = X[ filter_test , :]
if feature_subset != None:
    cols = [i for i, j in enumerate(continuous_features + discrete_features_dummies)
```

```

X_train = X_train[ : , cols ]
X_test = X_test[ : , cols ]

y_train = y[ filter_train ]
y_test = y[ filter_test ]

# Scale the variables
scaler = preprocessing.MinMaxScaler()

X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# return training and testing data
out = {'X_train':X_train, 'y_train':y_train, 'train_set':filter_train,
       'X_test':X_test, 'y_test':y_test, 'test_set':filter_test}

return out

def fit_classification(model, data_dict,
                      cv_parameters = {},
                      model_name = None,
                      random_state = default_seed,
                      output_to_file = True,
                      print_to_screen = True):
    ...
    This function will fit a classification model to data and print various evaluation
    measures. It expects the following parameters
    - model: an sklearn model object
    - data_dict: the dictionary containing both training and testing data;
                  returned by the prepare_data function
    - cv_parameters: a dictionary of parameters that should be optimized
                     over using cross-validation. Specifically, each named
                     entry in the dictionary should correspond to a parameter,
                     and each element should be a list containing the values
                     to optimize over
    - model_name: the name of the model being fit, for printouts
    - random_state: the random seed to use
    - output_to_file: if the results will be saved to the output file
    - print_to_screen: if the results will be printed on screen

    If the model provided does not have a predict_proba function, we will
    simply print accuracy diagnostics and return.

    If the model provided does have a predict_proba function, we first
    figure out the optimal threshold that maximizes the accuracy and
    print out accuracy diagnostics. We then print an ROC curve, sensitivity/
    specificity curve, and calibration curve.

    This function returns a dictionary with the following entries
    - model: the best fitted model
    - y_pred: predictions for the test set
    - y_pred_probs: probability predictions for the test set, if the model
                    supports them
    - y_pred_score: prediction scores for the test set, if the model does not
                    output probabilities.
    ...

    np.random.seed(random_state)

    # -----
    # Step 1 - Load the data
    # -----
    X_train = data_dict['X_train']
    y_train = data_dict['y_train']

    X_test = data_dict['X_test']
    y_test = data_dict['y_test']

```

```

filter_train = data_dict['train_set']

# -----
# Step 2 - Fit the model
# -----

cv_model = GridSearchCV(model, cv_parameters)

start_time = time.time()
cv_model.fit(X_train, y_train)
end_time = time.time()

best_model = cv_model.best_estimator_

if print_to_screen:
    if model_name != None:
        print("=====")
        print(" Model: " + model_name)
        print("=====")

    print("Fit time: " + str(round(end_time - start_time, 2)) + " seconds")
    print("Optimal parameters:")
    print(cv_model.best_params_)
    print("")

# -----
# Step 3 - Evaluate the model
# -----

# If possible, make probability predictions
try:
    y_pred_probs = best_model.predict_proba(X_test)[: ,1]
    fpr, tpr, thresholds = roc_curve(y_test, y_pred_probs)

    probs_predicted = True
except:
    probs_predicted = False

# Make predictions; if we were able to find probabilities, use
# the threshold that maximizes the accuracy in the training set.
# If not, just use the learner's predict function
if probs_predicted:
    y_train_pred_probs = best_model.predict_proba(X_train)[: ,1]
    fpr_train, tpr_train, thresholds_train = roc_curve(y_train, y_train_pred_probs)

    true_pos_train = tpr_train*(y_train.sum())
    true_neg_train = (1 - fpr_train) *(1-y_train).sum()

    best_threshold_index = np.argmax(true_pos_train + true_neg_train)
    best_threshold = 1 if best_threshold_index == 0 else thresholds_train[ best_threshold_index]

    if print_to_screen:
        print("Accuracy-maximizing threshold was: " + str(best_threshold))

    y_pred = (y_pred_probs > best_threshold)
else:
    y_pred = best_model.predict(X_test)

if print_to_screen:
    print("Accuracy: ", accuracy_score(y_test, y_pred))
    print(classification_report(y_test, y_pred, target_names = ['No default', 'Default'], ))

if print_to_screen:
    if probs_predicted:
        plt.figure(figsize = (13, 4.5))
        plt.subplot(2, 2, 1)

        plt.title("ROC Curve (AUC = %0.2f)"% roc_auc_score(y_test, y_pred_probs))

```

```

plt.plot(fpr, tpr, 'b')
plt.plot([0,1],[0,1],'r--')
plt.xlim([0,1]); plt.ylim([0,1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')

plt.subplot(2, 2, 3)

plt.plot(thresholds, tpr, 'b', label = 'Sensitivity')
plt.plot(thresholds, 1 - fpr, 'r', label = 'Specificity')
plt.legend(loc = 'lower right')
plt.xlim([0,1]); plt.ylim([0,1])
plt.xlabel('Threshold')

plt.subplot(2, 2, 2)

fp_0, mpv_0 = calibration_curve(y_test, y_pred_probs, n_bins = 10)
plt.plot([0,1], [0,1], 'k:', label='Perfectly calibrated')
plt.plot(mpv_0, fp_0, 's-')
plt.ylabel('Fraction of Positives')
plt.xlim([0,1]); plt.ylim([0,1])
plt.legend(loc = 'upper left')

plt.subplot(2, 2, 4)
plt.hist(y_pred_probs, range=(0, 1), bins=10, histtype="step", lw=2)
plt.xlim([0,1]); plt.ylim([0,20000])
plt.xlabel('Mean Predicted Probability')
plt.ylabel('Count')

#plt.tight_layout()
plt.show()

# Additional Score Check
if probs_predicted:
    y_train_score = y_train_pred_probs
else:
    y_train_score = best_model.decision_function(X_train)

tau, p_value = kendalltau(y_train_score, data.grade[filter_train])
if print_to_screen:
    print("")
    print("Similarity to LC grade ranking: ", tau)

if probs_predicted:
    brier_score = brier_score_loss(y_test, y_pred_probs)
    if print_to_screen:
        print("Brier score:", brier_score)

# Return the model predictions, and the
# test set
# -----
out = {'model':best_model, 'y_pred_labels':y_pred, 'accuracy':accuracy_score(y_test, y_pred_labels)}

if probs_predicted:
    out.update({'y_pred_probs':y_pred_probs})
else:
    y_pred_score = best_model.decision_function(X_test)
    out.update({'y_pred_score':y_pred_score})

# Output results to file
# -----
if probs_predicted and output_to_file:
    # Check whether any of the CV parameters are on the edge of
    # the search space
    opt_params_on_edge = find_opt_params_on_edge(cv_model)
    dump_to_output(model_name + "::search_on_edge", opt_params_on_edge)
    if print_to_screen:
        print("Were parameters on edge? : " + str(opt_params_on_edge))

# Find out how different the scores are for the different values

```

```

# tested for by cross-validation. If they're not too different, then
# even if the parameters are off the edge of the search grid, we should
# be ok
score_variation = find_score_variation(cv_model)
dump_to_output(model_name + "::score_variation", score_variation)
if print_to_screen:
    print("Score variations around CV search grid : " + str(score_variation))

# Print out all the scores
dump_to_output(model_name + "::all_cv_scores", str(cv_model.cv_results_['mean_test_score']))
if print_to_screen:
    print( str(cv_model.cv_results_['mean_test_score']) )

# Dump the AUC to file
dump_to_output(model_name + "::roc_auc", roc_auc_score(y_test, y_pred_probs) )

return out

```

### ▼ 3.1.1. Train and Test different machine learning classification models

The machine learning models listed in the following are just our suggestions. You are free to try any other models that you would like to experiment with.

```

## define your set of features to use in different models

# your_features = ['grade', 'purpose', 'term', 'emp_length', 'fico_range_high', 'revol_util',
your_features = list(discrete_features + continuous_features)

# prepare the train, test data for training models
data_dict = prepare_data(feature_subset = your_features)

all_features = pd.Series(continuous_features + discrete_features_dummies)
idx = [i for i, j in enumerate(continuous_features + discrete_features_dummies)
        if j.split("::")[0] in your_features]

selected_features = all_features[idx]
selected_features.reset_index(drop=True, inplace=True)

```

### ▼ Naive Bayes

```

## Train and test a naive bayes classifier

gnb = GaussianNB()
gnb = fit_classification(gnb, data_dict, model_name='gnb')

```





```
Model: gnb
```

```
Fit time: 0.2 seconds
```



```
Optimal parameters:
```

```
{}
```

```
Accuracy-maximizing threshold was: 1
```

```
Accuracy: 0.8035
```

	precision	recall	f1-score	support
No default	0.8035	1.0000	0.8910	16070
Default	0.0000	0.0000	0.0000	3930
micro avg	0.8035	0.8035	0.8035	20000
macro avg	0.4017	0.5000	0.4455	20000
weighted avg	0.6456	0.8035	0.7160	20000



## ▼ $l_1$ regularized logistic regression

```
## Train and test a  $l_1$  regularized logistic regression classifier
```

```
l1_logistic = LogisticRegression(penalty='l1')
```

```
cv_parameters = {'solver': ['liblinear', 'saga'], 'C': [.1, .5, 1]}
```

```
l1_logistic = fit_classification(l1_logistic, data_dict, cv_parameters = cv_parameters, model_
```



```
Model: Logisitic Regression
```

```
Fit time: 52.43 seconds
```



```
Optimal parameters:
```

```
{'C': 1, 'solver': 'saga'}
```

```
Accuracy-maximizing threshold was: 0.488192032703962
```

```
Accuracy: 0.80365
```

	precision	recall	f1-score	support
No default	0.8111	0.9851	0.8897	16070
Default	0.5031	0.0618	0.1101	3930
micro avg	0.8036	0.8036	0.8036	20000
macro avg	0.6571	0.5234	0.4999	20000
weighted avg	0.7506	0.8036	0.7365	20000



## ▼ $l_2$ regularized logistic regression

```
## Train and test a  $l_1$  regularized logistic regression classifier
```

```
l2_logistic = LogisticRegression(penalty='l2')
```

```
cv_parameters = {'solver': ['liblinear', 'saga', 'sag', 'newton-cg', 'lbfgs'], 'C': [.1, 1, 10]}
```

```
l2_logistic = fit_classification(l2_logistic, data_dict, cv_parameters = cv_parameters, model_
```



```
=====
Model: Logisitic Regression - L2
=====
```

```
Fit time: 33.31 seconds
```


```
Optimal parameters:
```

```
{'C': 1, 'solver': 'lbfgs'}
```

```
Accuracy-maximizing threshold was: 0.4904474805677525
```

```
Accuracy: 0.8036
```

	precision	recall	f1-score	support
No default	0.8109	0.9853	0.8897	16070
Default	0.5021	0.0606	0.1081	3930
micro avg	0.8036	0.8036	0.8036	20000
macro avg	0.6565	0.5229	0.4989	20000
weighted avg	0.7502	0.8036	0.7361	20000



```
## plot top 3 features with the most positive (and negative) weights
```

```
top_and_bottom_idx = list(np.argsort(l2_logistic['model'].coef_[0,:3]) + list(np.argsort(l2_
```


```
bplot = pd.Series(l2_logistic['model'].coef_[0,top_and_bottom_idx])
```

```
xticks = selected_features[top_and_bottom_idx]
```

```
p1 = bplot.plot(kind='bar',rot=-30,ylim=(-5,10))
```

```
p1.set_xticklabels(xticks)
```

```
plt.show()
```



## ▼ Decision tree

```
## Train and test a decision tree classifier
```

```
decision_tree = DecisionTreeClassifier()
```

```
cv_parameters = {'max_depth':[None, 3, 10, 15, 20, 100]}
```

```
decision_tree = fit_classification(decision_tree, data_dict, cv_parameters=cv_parameters, mod
```

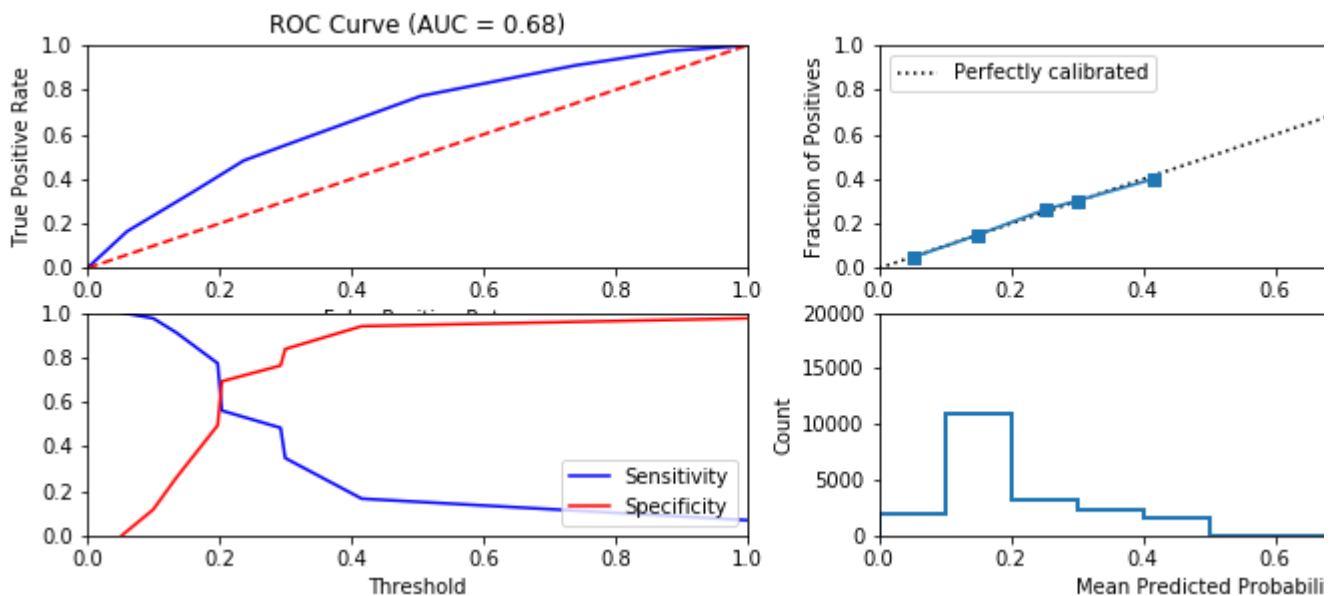


```
=====
Model: Decision Tree Classifier
=====
Fit time: 5.89 seconds
Optimal parameters:
{'max_depth': 3}
```

Accuracy-maximizing threshold was: 1

Accuracy: 0.8035

	precision	recall	f1-score	support
No default	0.8035	1.0000	0.8910	16070
Default	0.0000	0.0000	0.0000	3930
micro avg	0.8035	0.8035	0.8035	20000
macro avg	0.4017	0.5000	0.4455	20000
weighted avg	0.6456	0.8035	0.7160	20000



Similarity to LC grade ranking: 0.7787166588431091

Brier score: 0.14770363845538237

Were parameters on edge? : False

Score variations around CV search grid : 13.276648693488182

[0.69806667 0.80366667 0.7769 0.73946667 0.71466667 0.69696667]

## ▼ Random forest

```
## Train and test a random forest classifier
```

```
random_forest = RandomForestClassifier()
cv_parameters = {'n_estimators':[2, 5, 10, 20, 50], 'max_depth':[None, 2, 3, 5, 10]}
```

```
random_forest = fit_classification(random_forest, data_dict, cv_parameters=cv_parameters, mod
```



```
Model: Random Forest Classifier
```

```
Fit time: 30.43 seconds
```

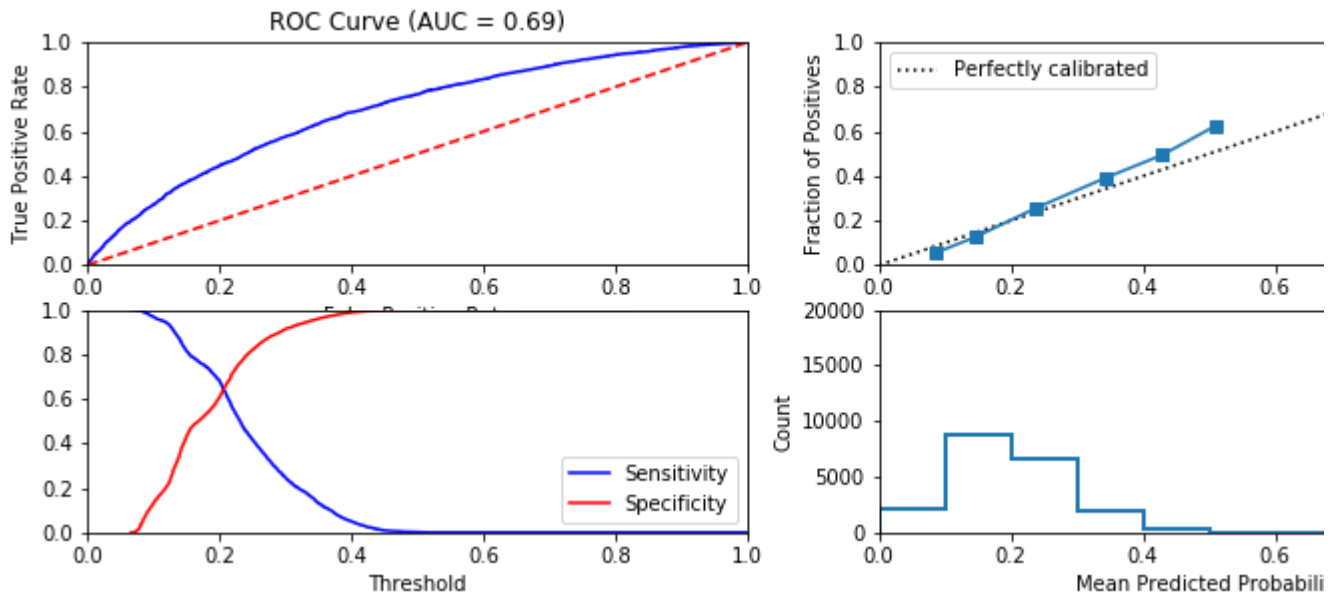
```
Optimal parameters:
```

```
{'max_depth': 5, 'n_estimators': 10}
```

```
Accuracy-maximizing threshold was: 0.3799892945891109
```

```
Accuracy: 0.80095
```

	precision	recall	f1-score	support
No default	0.8124	0.9781	0.8876	16070
Default	0.4609	0.0766	0.1314	3930
micro avg	0.8010	0.8010	0.8010	20000
macro avg	0.6367	0.5273	0.5095	20000
weighted avg	0.7434	0.8010	0.7390	20000



```
Similarity to LC grade ranking: 0.7620794934264274
```

```
Brier score: 0.14664377872096301
```

```
Were parameters on edge? : False
```


```
Score variations around CV search grid : 4.574296022892209
```

```
[0.78053333 0.767      0.79433333 0.79836667 0.8021      0.80363333
 0.8037      0.8037      0.8037      0.8037      0.80363333 0.8037
 0.8037      0.8037      0.8037      0.8028      0.80333333 0.80376667
 0.8037      0.8037      0.78546667 0.79993333 0.8023      0.80333333
 0.8033      ]
```

```
## Plot top 6 most significant features
```

```
top_idx = list(np.argsort(random_forest['model'].feature_importances_)[-6:])
bplot = pd.Series(random_forest['model'].feature_importances_[top_idx])
xticks = selected_features[top_idx]
p2 = bplot.plot(kind='bar',rot=-30,ylim=(0,0.2))
p2.set_xticklabels(xticks)
plt.show()
```





## ▼ Multi-layer perceptron

## Train and test a multi-layer perceptron classifier

```
mlp = MLPClassifier()
cv_parameters = {'activation':['logistic', 'tanh', 'relu'], 'hidden_layer_sizes':[(16, 8, 4,)]
mlp = fit_classification(mlp, data_dict, cv_parameters=cv_parameters, model_name='MLP Classif:
```



```
=====
Model: MLP Classifier
=====
Fit time: 437.31 seconds
Optimal parameters:
{'activation': 'logistic', 'hidden_layer_sizes': (16,)}

Accuracy-maximizing threshold was: 0.4773280825758916
Accuracy: 0.80305
```

	precision	recall	f1-score	support
No default	0.8104	0.9855	0.8894	16070
Default	0.4902	0.0570	0.1021	3930

## Train and Test logistic regression model with features derived by LendingClub

ROC Curve (AUC = 0.71)

```
## Find a lendingClub-defined feature and train a l1-regularized logistic regression model on
a_lendingclub_feature = ['grade']

data_dict = prepare_data(feature_subset = a_lendingclub_feature)
lc1_only_logistic = LogisticRegression(penalty='l1')
cv_parameters = {'solver': ['liblinear', 'saga'], 'C': [.1, .5, 1]}

lc1_only_logistic = fit_classification(lc1_only_logistic, data_dict, cv_parameters=cv_parameters)
```



```
=====
Model: Logistic Regression - L1 - Single Feature
=====
```

```
Fit time: 2.26 seconds
```

```
Optimal parameters:
```

```
{'C': 0.1, 'solver': 'liblinear'}
```

```
Accuracy-maximizing threshold was: 0.4502513698105868
```

```
Accuracy: 0.8035
```

	precision	recall	f1-score	support
No default	0.8035	1.0000	0.8910	16070
Default	0.0000	0.0000	0.0000	3930

```
## train a l2-regularized logistic regression model on data with only that feature
```

```
lc2_only_logistic = LogisticRegression(penalty='l2')
```

```
cv_parameters = {'solver': ['liblinear', 'saga', 'sag', 'newton-cg', 'lbfgs'], 'C': [.1, 1, 10]}
```

```
lc2_only_logistic = fit_classification(lc2_only_logistic, data_dict, cv_parameters=cv_parameters)
```





```

=====
Model: Logistic Regression - L2 - Single Feature
=====
Fit time: 4.91 seconds
Optimal parameters:
{'C': 0.1, 'solver': 'liblinear'}

Accuracy-maximizing threshold was: 0.4014245746842859
Accuracv: 0.8035

```

## Train and test all the models you have tried previously after removing features derived by LendingClub

```

# excluding 'grade' and 'int_rate'
non_lending_club_features = list(set(your_features) - set(['grade', 'int_rate']))

data_dict = prepare_data(feature_subset = non_lending_club_features)

gnb_nonlc = GaussianNB()
l1_logistic_nonlc = LogisticRegression(penalty='l1', solver='saga', C=1, random_state=0)
l2_logistic_nonlc = LogisticRegression(penalty='l2', solver='lbfgs', C=1, random_state=0)
dtc_nonlc = DecisionTreeClassifier(max_depth=3, random_state=0)
rfc_nonlc = RandomForestClassifier(n_estimators=20, max_depth=10, random_state=0)
mlp_nonlc = MLPClassifier(activation='logistic', hidden_layer_sizes=(16,), random_state=0)

model_dict = {'Naive Bayes': gnb_nonlc, 'L1 Logistic': l1_logistic_nonlc, 'L2 Logistic': l2_logistic_nonlc}

model_scores = {}

for n, m in model_dict.items():

    acc_scores = []
    roc_scores = []

    for i in range(100):
        data_dict = prepare_data(feature_subset = non_lending_club_features, random_state=i)
        model = fit_classification(m, data_dict, model_name=n, print_to_screen=False)
        acc_scores.append(model['accuracy'])
        roc_scores.append(model['ROC'])

    avg_acc = np.mean(acc_scores)
    std_acc = np.std(acc_scores)
    avg_roc = np.mean(roc_scores)
    std_roc = np.std(roc_scores)

    model_scores[n] = [avg_acc, std_acc, avg_roc, std_roc]

print("=====")
print(n)
print("=====")
print("Average Accuracy: ", round(avg_acc, 5), "+/-", round(std_acc, 5))
print("Average ROC: ", round(avg_roc, 5), "+/-", round(std_roc, 5))
print()

```



```

=====
Naive Bayes
=====
Average Accuracy:  0.80209 +/- 0.00277
Average ROC:  0.64812 +/- 0.00523

=====
L1 Logistic
=====
Average Accuracy:  0.80288 +/- 0.00286
Average ROC:  0.69093 +/- 0.00425

=====
L2 Logistic
=====
Average Accuracy:  0.80296 +/- 0.00275
Average ROC:  0.69067 +/- 0.00424

=====
Decision Tree
=====
Average Accuracy:  0.80212 +/- 0.00276
Average ROC:  0.64813 +/- 0.00498

=====
Random Forest
=====
Average Accuracy:  0.78869 +/- 0.00431
Average ROC:  0.67779 +/- 0.00429

=====
MLP
=====
Average Accuracy:  0.80296 +/- 0.00285

```

```

data_dict = prepare_data(feature_subset = non_lending_club_features)
our_model = LogisticRegression(penalty='l1', solver='saga', C=1)
our_model_fit = fit_classification(our_model, data_dict, model_name='Our Model: Non-LC Features')

```



```
=====
Model: Our Model: Non-LC Features
=====
```

```
Fit time: 9.52 seconds
```



```
Optimal parameters:
```

```
{}
```

```
Accuracy-maximizing threshold was: 0.5024496981358255
```

```
Accuracy: 0.80435
```

	precision	recall	f1-score	support
No default	0.8102	0.9880	0.8903	16070
Default	0.5211	0.0534	0.0969	3930
micro avg	0.8044	0.8044	0.8044	20000
macro avg	0.6656	0.5207	0.4936	20000
weighted avg	0.7534	0.8044	0.7344	20000



## ▼ Time stability test of YOURMODEL

```
0.0 1          15000 1

## Define the time window of your train and test data
## First run with 2010 training data
start_date_train = datetime.datetime.strptime( 'Jan-2010', "%b-%Y").date()
end_date_train = datetime.datetime.strptime( 'Dec-2010', "%b-%Y").date()
start_date_test = datetime.datetime.strptime( 'Jan-2017', "%b-%Y").date()
end_date_test = datetime.datetime.strptime( 'Dec-2017', "%b-%Y").date()

data_dict_test = prepare_data(date_range_train = (start_date_train, end_date_train),
                              date_range_test = (start_date_test, end_date_test),
                              n_samples_train = 7000, n_samples_test = 7000, feature_subset = your.

## Train and test YOURMODEL using this data
our_model = LogisticRegression(penalty='l1')
cv_parameters = {'solver': ['liblinear', 'saga'], 'C': [.1, .5, 1]}

our_model_fit = fit_classification(our_model, data_dict_test, cv_parameters=cv_parameters, mo
```



```
=====
Model: Our Model : Logistic Regression
=====
```

```
Fit time: 2.57 seconds
```

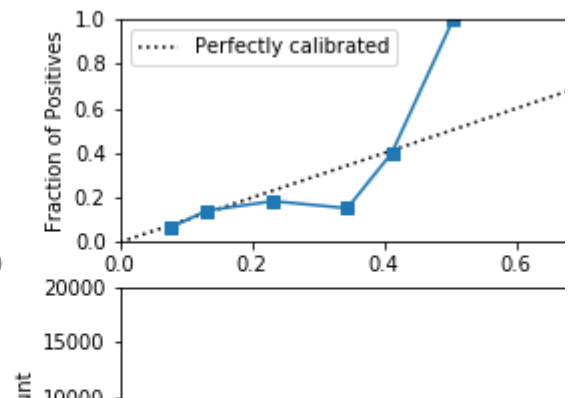

```
Optimal parameters:
```

```
{'C': 0.1, 'solver': 'liblinear'}
```

```
Accuracy-maximizing threshold was: 1
```

```
Accuracy: 0.8804285714285714
```

	precision	recall	f1-score	support
No default	0.8804	1.0000	0.9364	6163
Default	0.0000	0.0000	0.0000	837
micro avg	0.8804	0.8804	0.8804	7000
macro avg	0.4402	0.5000	0.4682	7000
weighted avg	0.7752	0.8804	0.8244	7000



```
## Run with 2016 training data
```

```
start_date_train = datetime.datetime.strptime('Jan-2016', "%b-%Y").date()
```

```
end_date_train = datetime.datetime.strptime('Dec-2016', "%b-%Y").date()
```

```
start_date_test = datetime.datetime.strptime('Jan-2017', "%b-%Y").date()
```

```
end_date_test = datetime.datetime.strptime('Dec-2017', "%b-%Y").date()
```

```
data_dict_test = prepare_data(date_range_train = (start_date_train, end_date_train),
                              date_range_test = (start_date_test, end_date_test),
                              n_samples_train = 7000, n_samples_test = 7000, feature_subset = your_
```

```
## Train and test YOURMODEL using this data
```

```
our_model = LogisticRegression(penalty='l1')
```

```
cv_parameters = {'solver': ['liblinear', 'saga'], 'C': [.1, .5, 1]}
```


```
our_model_fit = fit_classification(our_model, data_dict_test, cv_parameters=cv_parameters, mo
```



```
=====
Model: Our Model 2016: Logistic Regression
=====
Fit time: 4.25 seconds
Optimal parameters:
{'C': 0.1, 'solver': 'liblinear'}

Accuracy-maximizing threshold was: 0.48412672110776184
Accuracy: 0.8575714285714285
```

	precision	recall	f1-score	support
No default	0.8924	0.9539	0.9221	6188
Default	0.2597	0.1232	0.1671	812
micro avg	0.8576	0.8576	0.8576	7000
macro avg	0.5761	0.5385	0.5446	7000
weighted avg	0.8190	0.8576	0.8345	7000



## ▼ Train and test YOURMODEL on the original data

Pre-pickled data from Phase 2 to capture data in a state before cleaning and scaling, but after necessary type casting to ensure model runs. All "dirty" data is captured below as pre\_data. Added one-hot-encoding and converted 'loan\_status' to a binary 'outcome' value.

```
import pickle
import numpy as np
pre_data, pre_discrete_features, pre_continuous_features = pickle.load( open( "/gdrive/My Drive/

# in this step, we remove features with too many unique values (i.e. 'title' which had over 40
import pandas as pd

pre_discrete_cols = []
for f in pre_discrete_features:
    if pre_data[f].nunique() < 25:
        pre_discrete_cols.append(f)

# necessary to one-hot-encode discrete features so that the model runs
X_continuous_pre = pre_data[list(pre_continuous_features)].values
```

```
X_discrete_pre = pd.get_dummies(pre_data[pre_discrete_cols], dummy_na = True, prefix_sep = ":")
X_discrete_pre = X_discrete_pre.values

pre_X = np.concatenate( (X_continuous_pre, X_discrete_pre), axis = 1 )
pre_y = pre_data.outcome.values

pre_X_train, pre_X_test, pre_y_train, pre_y_test = train_test_split(pre_X, pre_y, test_size=.
```

```
from sklearn.linear_model import LogisticRegression
```

```
## Train and test YOURMODEL using this data
our_model = LogisticRegression(penalty='l1')
our_model.fit(pre_X_train, pre_y_train)
our_pred = our_model.predict_proba(pre_X_test)
```

```
our_pred2 = our_model.predict(pre_X_test).
```

```
roc_auc_score(pre_y_test, our_pred[:,1]).
```

```
↳ 0.9836979315475219
```

```
accuracy_score(pre_y_test, our_pred2).
```

```
↳ 0.9679089751569834
```

## ▼ Test regression models

```
def fit_regression(model, data_dict,
                  cv_parameters = {},
                  separate = False,
                  model_name = None,
                  random_state = default_seed,
                  output_to_file = True,
                  print_to_screen = True):
    ...
```

This function will fit a regression model to data and print various evaluation measures. It expects the following parameters

- model: an sklearn model object
- data\_dict: the dictionary containing both training and testing data; returned by the prepare\_data function
- separate: a Boolean variable indicating whether we fit models for defaulted and non-defaulted loans separately
- cv\_parameters: a dictionary of parameters that should be optimized over using cross-validation. Specifically, each named entry in the dictionary should correspond to a parameter, and each element should be a list containing the values to optimize over
- model\_name: the name of the model being fit, for printouts
- random\_state: the random seed to use
- output\_to\_file: if the results will be saved to the output file
- print\_to\_screen: if the results will be printed on screen

This function returns a dictionary FOR EACH RETURN DEFINITION with the following entries

- model: the best fitted model
- predicted\_return: prediction result based on the test set
- predicted\_regular\_return: prediction result for non-defaulted loans (valid if separate == False)
- predicted\_default\_return: prediction result for defaulted loans (valid if separate == True)

```

... - r2_scores: the testing r2_score(s) for the best fitted model
...

np.random.seed(random_state)

# -----
# Step 1 - Load the data
# -----

col_list = ['ret_PESS', 'ret_OPT', 'ret_INTa', 'ret_INTb']

X_train = data_dict['X_train']
filter_train = data_dict['train_set']

X_test = data_dict['X_test']
filter_test = data_dict['test_set']
out = {}

for ret_col in col_list:

    y_train = data.loc[filter_train, ret_col].as_matrix()
    y_test = data.loc[filter_test, ret_col].as_matrix()

    # -----
    # Step 2 - Fit the model
    # -----

    if separate:
        outcome_train = data.loc[filter_train, 'outcome']
        outcome_test = data.loc[filter_test, 'outcome']

        # Train two separate regressors for defaulted and non-defaulted loans
        X_train_0 = X_train[outcome_train == False]
        y_train_0 = y_train[outcome_train == False]
        X_test_0 = X_test[outcome_test == False]
        y_test_0 = y_test[outcome_test == False]

        X_train_1 = X_train[outcome_train == True]
        y_train_1 = y_train[outcome_train == True]
        X_test_1 = X_test[outcome_test == True]
        y_test_1 = y_test[outcome_test == True]

        cv_model_0 = GridSearchCV(model, cv_parameters, scoring='r2')
        cv_model_1 = GridSearchCV(model, cv_parameters, scoring='r2')

        start_time = time.time()
        cv_model_0.fit(X_train_0, y_train_0)
        cv_model_1.fit(X_train_1, y_train_1)
        end_time = time.time()

        best_model_0 = cv_model_0.best_estimator_
        best_model_1 = cv_model_1.best_estimator_

    if print_to_screen:

        if model_name != None:
            print("=====")
            print(" Model: " + model_name + " Return column: " + ret_col)
            print("=====")

            print("Fit time: " + str(round(end_time - start_time, 2)) + " seconds")
            print("Optimal parameters:")
            print("model_0:", cv_model_0.best_params_, "model_1:", cv_model_1.best_params_)

        predicted_regular_return = best_model_0.predict(X_test)
        predicted_default_return = best_model_1.predict(X_test)

    if print_to_screen:
        print("")
        print("Testing r2 scores:")

```

```

# Here we use different testing set to report the performance
test_scores = {'model_0':r2_score(y_test_0,best_model_0.predict(X_test_0)),
               'model_1':r2_score(y_test_1,best_model_1.predict(X_test_1))}
if print_to_screen:
    print("model_0:", test_scores['model_0'])
    print("model_1:", test_scores['model_1'])

cv_objects = {'model_0':cv_model_0, 'model_1':cv_model_1}
out[ret_col] = { 'model_0':best_model_0, 'model_1':best_model_1, 'predicted_regul:
                'predicted_default_return':predicted_default_return, 'r2_scores':

else:
    cv_model = GridSearchCV(model, cv_parameters, scoring='r2')

    start_time = time.time()
    cv_model.fit(X_train, y_train)
    end_time = time.time()

    best_model = cv_model.best_estimator_

    if print_to_screen:
        if model_name != None:
            print("=====")
            print("  Model: " + model_name + "  Return column: " + ret_col)
            print("=====")

            print("Fit time: " + str(round(end_time - start_time, 2)) + " seconds")
            print("Optimal parameters:")
            print(cv_model.best_params_)

        predicted_return = best_model.predict(X_test)
        test_scores = {'model':r2_score(y_test,predicted_return)}
        if print_to_screen:
            print("")
            print("Testing r2 score:", test_scores['model'])

        cv_objects = {'model':cv_model}
        out[ret_col] = {'model':best_model, 'predicted_return':predicted_return, 'r2_score

# Output the results to a file
if output_to_file:
    for i in cv_objects:
        # Check whether any of the CV parameters are on the edge of
        # the search space
        opt_params_on_edge = find_opt_params_on_edge(cv_objects[i])
        dump_to_output(model_name + ":@" + ret_col + ":@"search_on_edge", opt_params_on
        if print_to_screen:
            print("Were parameters on edge (" + i + ") : " + str(opt_params_on_edge))

        # Find out how different the scores are for the different values
        # tested for by cross-validation. If they're not too different, then
        # even if the parameters are off the edge of the search grid, we should
        # be ok
        score_variation = find_score_variation(cv_objects[i])
        dump_to_output(model_name + ":@" + ret_col + ":@"score_variation", score_varia
        if print_to_screen:
            print("Score variations around CV search grid (" + i + ") : " + str(score

        # Print out all the scores
        dump_to_output(model_name + ":@"all_cv_scores", str(cv_objects[i].cv_results_
        if print_to_screen:
            print("All test scores : " + str(cv_objects[i].cv_results_['mean_test_sco

        # Dump the AUC to file
        dump_to_output( model_name + ":@" + ret_col + ":@"r2", test_scores[i] )

return out

```



## ▼ $l_1$ regularized linear regression

```
data_dict = prepare_data(feature_subset = your_features)
```

```
## First, trying l1 regularized linear regression with hyper-parameters
```

```
l1_linear = linear_model.LinearRegression()
```

```
reg_lasso = fit_regression(l1_linear, data_dict, model_name="Lasso Regression")
```



## ▼ $l_2$ regularized linear regressor

fit time: 0.33 seconds

```
## trying l2 regularized linear regression with hyper-parameters
```

```
l2_linear = linear_model.Ridge()  
cv_parameters = {'alpha': [.1, 1, 10, 100]}
```

```
reg_ridge = fit_regression(l2_linear, data_dict, cv_parameters=cv_parameters, model_name="Ridge")
```



## ▼ Multi-layer perceptron regressor

Fit time: 0.5 seconds

```
## trying multi-layer perceptron regression with hyper-parameters
```

```
mlr = MLPRegressor()
```

```
cv_parameters = {'activation':['logistic', 'tanh', 'relu'], 'hidden_layer_sizes':[(16, 8, 4),
```

```
reg_mlp = fit_regression(mlr, data_dict, cv_parameters=cv_parameters, model_name='MLP Regressor')
```



```
=====
Model: MLP Regressor Return column: ret_PESS
=====
Fit time: 44.68 seconds
```

### ▼ Random forest regressor

```
Testing r2 score: 0.023751170772329195
```

```
## trying random forest regression with hyper-parameters
```

```
rfr = RandomForestRegressor()
```

```
cv_parameters = {'n_estimators':[2, 5, 10, 20, 50], 'max_depth':[None, 2, 3, 5, 10]}
```

```
reg_rf = fit_regression(rfr, data_dict, cv_parameters=cv_parameters, model_name='Random Forest')
```



```

=====
Model: Random Forest Regressor Return column: ret_PESS
=====
Fit time: 159.64 seconds
Optimal parameters:
{'max_depth': 10, 'n_estimators': 50}

Testing r2 score: 0.044241948981436074
Were parameters on edge (model) : True
Score variations around CV search grid (model) : -519.4545116662637
All test scores : [-0.70952894 -0.35666261 -0.24288679 -0.18233118 -0.15182048 -0.1414
-0.13544519 -0.13770133 -0.14079087 -0.1380834 -0.14573572 -0.13800824
-0.13511484 -0.13283199 -0.13284499 -0.14777181 -0.13217382 -0.1294526
-0.1190723 -0.12145063 -0.23474641 -0.1628316 -0.13856809 -0.12325707
-0.11454093]
=====
Model: Random Forest Regressor Return column: ret_OPT
=====
Fit time: 155.15 seconds
Optimal parameters:
{'max_denth': 5, 'n_estimators': 50}

```

## ▼ Test investment strategies

Now we test several investment strategies using the learning models above

```

def test_investments(data_dict,
                     classifier = None,
                     regressor = None,
                     strategy = 'Random',
                     num_loans = 1000,
                     random_state = default_seed,
                     output_to_file = True):
    ...

```

This function tests a variety of investment methodologies and their returns. It will run its tests on the loans defined by the test\_set element of the data dictionary.

It is currently able to test four strategies

- random: invest in a random set of loans
- default-based: score each loan by probability of default, and only invest in the "safest" loans (i.e., those with the lowest probabilities of default)
- return-based: train a single regression model to predict the expected return of loans in the past. Then, for loans we could invest in, simply rank them by their expected returns and invest in that order.
- default-& return-based: train two regression models to predict the expected return of defaulted loans and non-defaulted loans in the training set. Then, for each potential loan we could invest in, predict the probability the loan will default, its return if it doesn't default and its return if it does. Then, calculate a weighted combination of the latter using the former to find a predicted return. Rank the loans by this expected return, and invest in that order

It expects the following parameters

- data\_dict: the dictionary containing both training and testing data; returned by the prepare\_data function
- classifier: a fitted model object which is returned by the fit\_classification function.
- regressor: a fitted model object which is returned by the fit\_regression function.
- strategy: the name of the strategy; one of the three listed above
- num\_loans: the number of loans to be included in the test portfolio
- num\_samples: the number of random samples used to compute average return ()

- random\_state: the random seed to use when selecting a subset of rows
- output\_to\_file: if the results will be saved to the output file

The function returns a dictionary FOR EACH RETURN DEFINITION with the following entries

- strategy: the name of the strategy
- average return: the return of the strategy based on the testing set
- test data: the updated Dataframe of testing data. Useful in the optimization section

```

np.random.seed(random_state)

# Retrieve the rows that were used to train and test the
# classification model
train_set = data_dict['train_set']
test_set = data_dict['test_set']

col_list = ['ret_PESS', 'ret_OPT', 'ret_INTa', 'ret_INTb']

# Create a dataframe for testing, including the score
data_test = data.loc[test_set,:]
out = {}

for ret_col in col_list:
    if strategy == 'Random':
        # Randomize the order of the rows in the dataframe
        data_test = data_test.sample(frac = 1).reset_index(drop = True)

        # Select num_loans to invest in
        pf_test = num_loans

        # Find the average return for these loans
        ret_test = np.mean(data_test[:pf_test][ret_col])

        # Return
        out[ret_col] = {'strategy':strategy, 'average return':ret_test}

        # Dump the strategy performance to file
        if output_to_file:
            dump_to_output(strategy + "," + ret_col + "::average return", ret_test )

        continue

    elif strategy == 'Return-based':

        colname = 'predicted_return_' + ret_col

        data_test[colname] = regressor[ret_col]['predicted_return']

        # Sort the loans by predicted return
        data_test = data_test.sort_values(by=colname, ascending = False).reset_index(drop

        ## Pick num_loans loans
        pf_test = num_loans

        ## Find their return
        ret_test = np.mean(data_test[:pf_test][ret_col])

        # Return
        out[ret_col] = {'strategy':strategy, 'average return':ret_test, 'test data':data_

        # Dump the strategy performance to file
        if output_to_file:
            dump_to_output(strategy + "," + ret_col + "::average return", ret_test )

        continue

    # Get the predicted scores, if the strategy is not Random or just Regression
    try:
        y_pred_score = classifier['y_pred_probs']

```

```

except:
    y_pred_score = classifier['y_pred_score']

data_test['score'] = y_pred_score

if strategy == 'Default-based':
    # Sort the test data by the score
    data_test = data_test.sort_values(by='score').reset_index(drop = True)

    ## Select num_loans to invest in
    pf_test = num_loans

    ## Find the average return for these loans
    ret_test = np.mean(data_test[:pf_test][ret_col])

    # Return
    out[ret_col] = {'strategy':strategy, 'average return':ret_test}

    # Dump the strategy performance to file
    if output_to_file:
        dump_to_output(strategy + "," + ret_col + "::average return", ret_test )

    continue

elif strategy == 'Default-return-based':

    # Load the predicted returns
    data_test['predicted_regular_return'] = regressor[ret_col]['predicted_regular_return']
    data_test['predicted_default_return'] = regressor[ret_col]['predicted_default_return']

    # Compute expectation
    colname = 'predicted_return_' + ret_col

    data_test[colname] = ( (1-data_test.score)*data_test.predicted_regular_return +
                          data_test.score*data_test.predicted_default_return )

    # Sort the loans by predicted return
    data_test = data_test.sort_values(by=colname, ascending = False).reset_index(drop

    ## Pick num_loans loans
    pf_test = num_loans

    ## Find their return
    ret_test = np.mean(data_test[:pf_test][ret_col])

    # Return
    out[ret_col] = {'strategy':strategy, 'average return':ret_test, 'test data':data_

    # Dump the strategy performance to file
    if output_to_file:
        dump_to_output(strategy + "," + ret_col + "::average return", ret_test )

    continue

else:
    return 'Not a valid strategy'

return out

## Test investment strategies using the best performing regressor

col_list = ['ret_PESS', 'ret_OPT', 'ret_INTa', 'ret_INTb']
test_strategy = 'Random'

print('strategy:',test_strategy)
strat_rand = test_investments(data_dict, regressor=reg_lasso, classifier=l1_logistic, strategy=

```

```
for ret_col in col_list:
    print(ret_col + ': ' + str(strat_rand[ret_col]['average return']))
```

```
↳ strategy: Random
   ret_PESS: 0.003643181247051625
   ret_OPT: 0.04185929462338522
   ret_INTa: 0.01952620210601705
   ret_INTb: 0.05513190207281548
```

```
test_strategy = 'Default-based'
```

```
print('strategy:', test_strategy)
strat_def = test_investments(data_dict, regressor=reg_lasso, classifier=l1_logistic, strategy:
```

```
for ret_col in col_list:
    print(ret_col + ': ' + str(strat_def[ret_col]['average return']))
```

```
↳ strategy: Default-based
   ret_PESS: 0.01996309742233598
   ret_OPT: 0.05098787777902832
   ret_INTa: 0.020961224426680993
   ret_INTb: 0.05481708207671865
```

```
test_strategy = 'Return-based'
```

```
print('strategy:', test_strategy)
strat_ret = test_investments(data_dict, regressor=reg_lasso, classifier=l1_logistic, strategy:
```

```
for ret_col in col_list:
    print(ret_col + ': ' + str(strat_ret[ret_col]['average return']))
```

```
↳ strategy: Return-based
   ret_PESS: 0.031883880168584254
   ret_OPT: 0.04350940291941102
   ret_INTa: 0.02192249818373398
   ret_INTb: 0.053310209289889686
```

```
test_strategy = 'Default-return-based'
```

```
## For the Default-return-based strategy we need to fit a new regressor with separate = True
l1_linear_sep = linear_model.LinearRegression()
reg_separate = fit_regression(l1_linear_sep, data_dict, model_name="Lasso Regression", separa
```

```
print('strategy:', test_strategy)
strat_defret = test_investments(data_dict, regressor=reg_separate, classifier=l1_logistic, st
```

```
for ret_col in col_list:
    print(ret_col + ': ' + str(strat_defret[ret_col]['average return']))
```

```
↳ strategy: Default-return-based
   ret_PESS: 0.03442041235561199
   ret_OPT: 0.044773690367724606
   ret_INTa: 0.02841656791507316
   ret_INTb: 0.05475944447130236
```



## ▼ Sensitivity test of portfolio size

```
col_list = ['ret_PESS', 'ret_OPT', 'ret_INTa', 'ret_INTb']
strategy_dict = {'Random':strat_rand, 'Default-based':strat_def, 'Return-based':strat_ret, 'D
model_dict_reg = {'Lasso': reg_lasso, 'Ridge': reg_ridge, 'MLP': reg_mlp, 'Random Forest': re

returns = {k:[] for k in col_list}

for n, s in strategy_dict.items():

    print("=====")
    print(n)
    print("=====")

    for i in range(100):

        data_dict = prepare_data(feature_subset = your_features, random_state=i)

        if (n == 'Default-return-based'):
            strat_current = test_investments(data_dict, regressor=reg_separate, classifier=l1_logis
        else:
            strat_current = test_investments(data_dict, regressor=reg_lasso, classifier=l1_logistic

        for ret_col in col_list:
            returns[ret_col].append(strat_current[ret_col]['average return'])

    for ret_col in col_list:
        avg_ret = np.mean(returns[ret_col])
        std_ret = np.std(returns[ret_col])
        print("Average Return: ", ret_col, ": ", round(avg_ret, 5), "+/-", round(std_ret,5))

    print()
```



```
=====
Random
=====
Average Return:  ret_PESS :  0.00465 +/- 0.00286
Average Return:  ret_OPT :  0.04568 +/- 0.00383
Average Return:  ret_INTa : 0.02147 +/- 0.0019
Average Return:  ret_INTb : 0.05659 +/- 0.00227
```

```
## Test the best-performing data-driven strategy on different portfolio sizes
```

```
result_sensitivity = []
```

```
test_strategy = 'Return-based'
```

```
## Vary the portfolio size from 1,000 to 10,000
```

```
for num_loans in list(range(1000,10000,1000)):
```

```
    reg_0 = test_investments(data_dict, regressor=reg_lasso, classifier=l1_logistic, strategy:
    result_sensitivity.append(reg_0['ret_PESS']['average return'])
```

```
result_sensitivity = np.array(result_sensitivity) * 100
```

```
sns.pointplot(np.array(list(range(1000,10000,1000))),result_sensitivity)
```

```
sns.despine()
```

```
plt.ylabel('Investment Return (%)',size = 14)
```

```
plt.xlabel('Portfolio Size',size = 14)
```

```
plt.show()
```

